

# Probabilistic Machine Learning with Julia

## Lecture 4 | Probabilistic Models for Time Series

Amirabbas Asadi

amir.asadi78@sharif.edu

2025

# Today's Topics

Today we will have an overview of parametric probabilistic models for time series data:

- Bayesian Neural Networks
- Autoregressive models and Dynamic Bayesian Networks
- Time series modeling with Bayesian Differential Equations
- Motivation for nonparametric models

# Poisson Regression

In the previous lecture, we discussed a simple form of Poisson Regression:

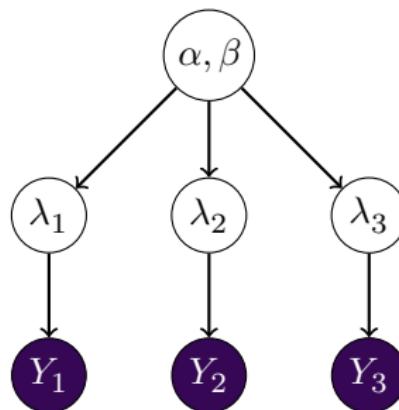
$$Y_\tau \sim \text{Poisson}(\lambda_\tau)$$

$$\log \lambda_\tau = \alpha\tau + \beta$$

$$\alpha, \beta \sim \text{Normal}(0, 1)$$

```
@model function poisson_linear_model(y)
    α ~ Normal(0.0, 1.0)
    β ~ Normal(0.0, 1.0)

    for τ ∈ 1:length(y)
        log_λ = α*τ + β
        y[τ] ~ Poisson(exp(log_λ))
    end
end
```

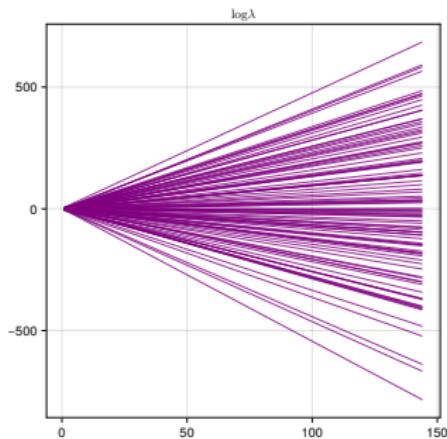


# Poisson Regression

$$p(\alpha, \beta) \xrightarrow{\text{Bayesian Inference}} p(\alpha, \beta|y)$$

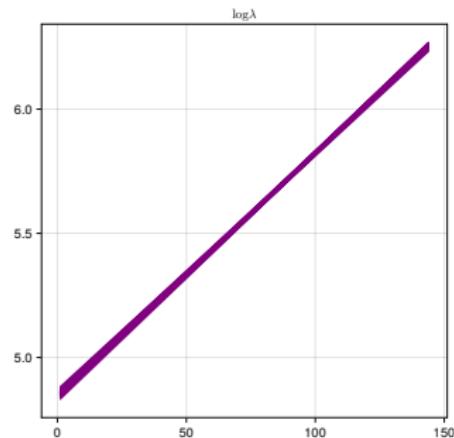
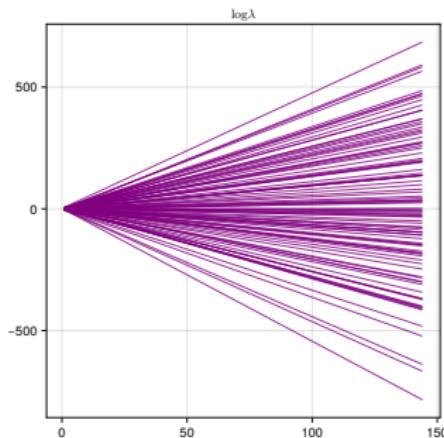
# Poisson Regression

$$p(\alpha, \beta) \xrightarrow{\text{Bayesian Inference}} p(\alpha, \beta|y)$$

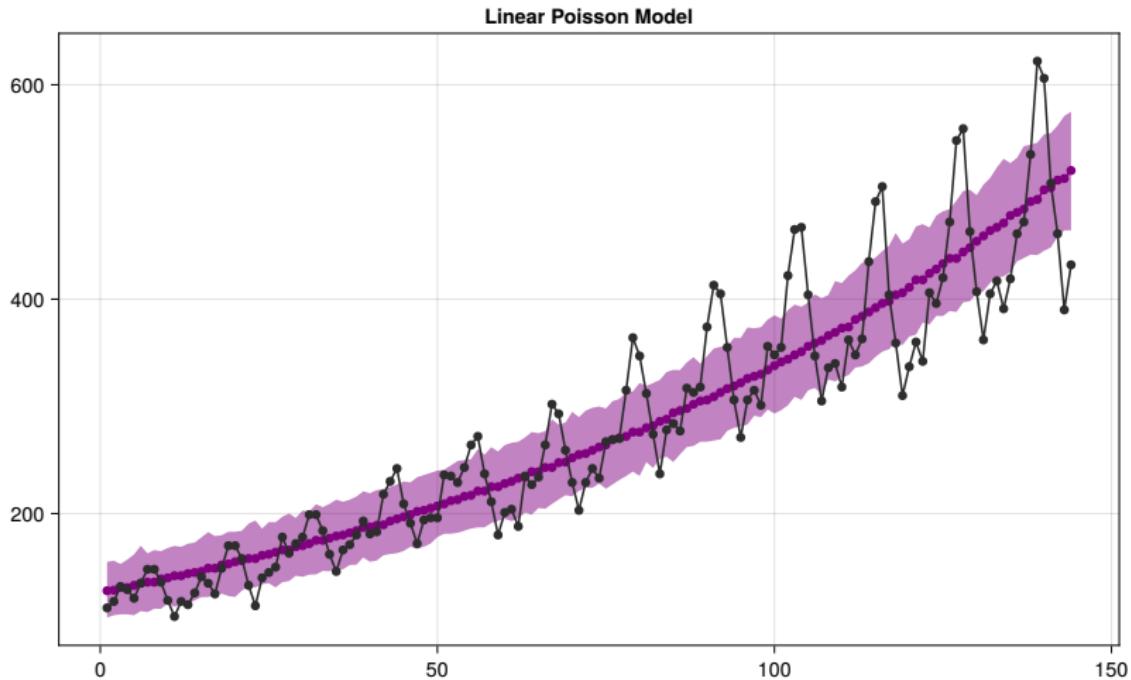


# Poisson Regression

$$p(\alpha, \beta) \xrightarrow{\text{Bayesian Inference}} p(\alpha, \beta|y)$$



# Poisson Regression



## Poisson Regression

It is easy to generalize this model by parametrizing  $\log \lambda_\tau$  using a nonlinear function:

$$\log \lambda_\tau = f(\cdot; w)$$

## Poisson Regression

It is easy to generalize this model by parametrizing  $\log \lambda_\tau$  using a nonlinear function:

$$\log \lambda_\tau = f(\cdot; w)$$

For example in the previous session, we added a nonlinear term for handling seasonality:

$$\log \lambda_\tau = \alpha\tau + \nu \cos(\omega\pi(\tau - \phi)) + \beta$$

## Poisson Regression

It is easy to generalize this model by parametrizing  $\log \lambda_\tau$  using a nonlinear function:

$$\log \lambda_\tau = f(\cdot; w)$$

For example in the previous session, we added a nonlinear term for handling seasonality:

$$\log \lambda_\tau = \alpha\tau + \nu \cos(\omega\pi(\tau - \phi)) + \beta$$

We can also use more expressive functions like a neural network if necessary.

Let's discuss **Bayesian Neural Network** first for a classification problem and then for our time series example.

# Bayesian Neural Network

Let  $f(x; \mathbf{w})$  be a neural network with three layers and sigmoid activation:

$$\mathbf{w} = [\mathbf{w}^{L1}, \mathbf{w}^{L2}, \mathbf{w}^{L3}]$$

$$f(x; \mathbf{w}) = \sigma(\sigma(\sigma(x\mathbf{w}^{L1})\mathbf{w}^{L2})\mathbf{w}^{L3})$$

# Bayesian Neural Network

Let  $f(x; \mathbf{w})$  be a neural network with three layers and sigmoid activation:

$$\mathbf{w} = [\mathbf{w}^{L1}, \mathbf{w}^{L2}, \mathbf{w}^{L3}]$$

$$f(x; \mathbf{w}) = \sigma(\sigma(\sigma(x\mathbf{w}^{L1})\mathbf{w}^{L2})\mathbf{w}^{L3})$$

We define a Multivariate Normal prior over  $\mathbf{w}$

$$\mathbf{w} \sim \text{Normal}(0, I)$$

# Bayesian Neural Network

```
function neural_network(X, L1, L2, L3, H_dim)
    H = X * reshape(L1, (2, H_dim))
    H = sigmoid.(H)

    H = H * reshape(L2, (H_dim, H_dim))
    H = sigmoid.(H)

    O = H * reshape(L3, (H_dim, 1))
    O = sigmoid.(O)

    return O
end
```

# Bayesian Neural Network

```
@model function bayesian_neural_network(X, y, σ, H_dim)
    n_L1 = 2 * H_dim
    n_L2 = H_dim * H_dim
    n_L3 = H_dim

    Σ(n) = Diagonal(abs2.(σ .* ones(n)))

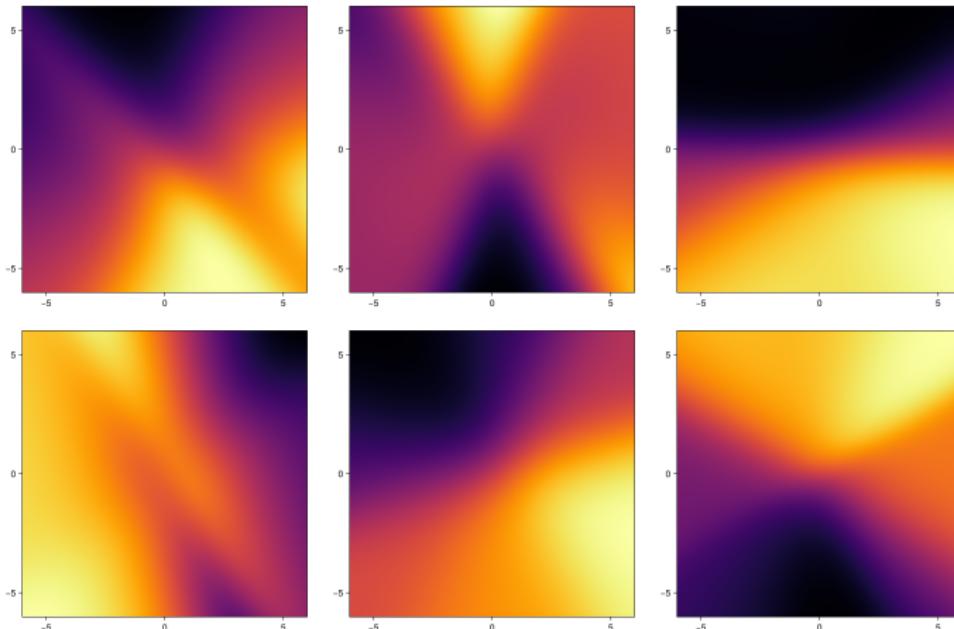
    L1 ~ MvNormal(zeros(n_L1), Σ(n_L1))
    L2 ~ MvNormal(zeros(n_L2), Σ(n_L2))
    L3 ~ MvNormal(zeros(n_L3), Σ(n_L3))

    O = neural_network(X, L1, L2, L3, H_dim)

    for i in eachindex(y)
        y[i] ~ Bernoulli(O[:, i])
    end
end
```

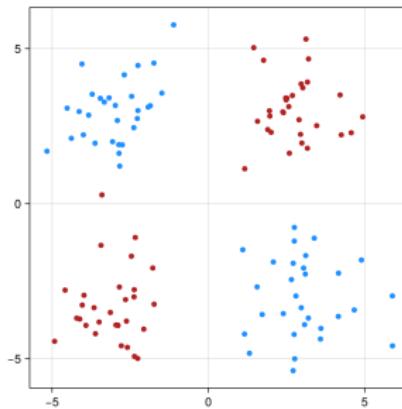
# Bayesian Neural Network

Some samples from  $p(\mathbf{w})$



# Bayesian Neural Network

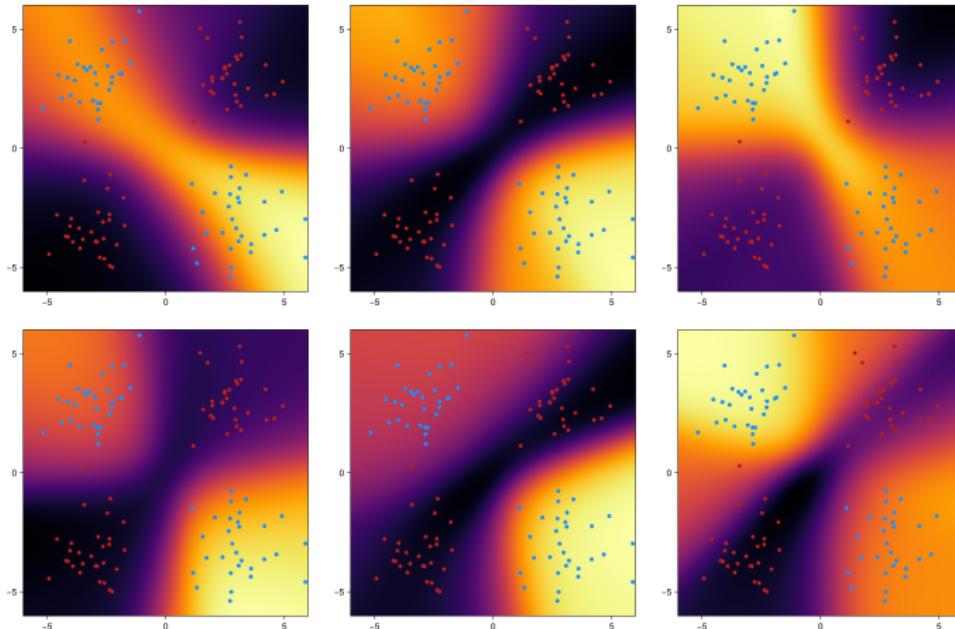
Now we update our belief about network weights using the below dataset



$$y \sim \text{Bernoulli}(f(x; \mathbf{w}))$$

# Bayesian Neural Network

Some samples from  $p(\mathbf{w}|y)$



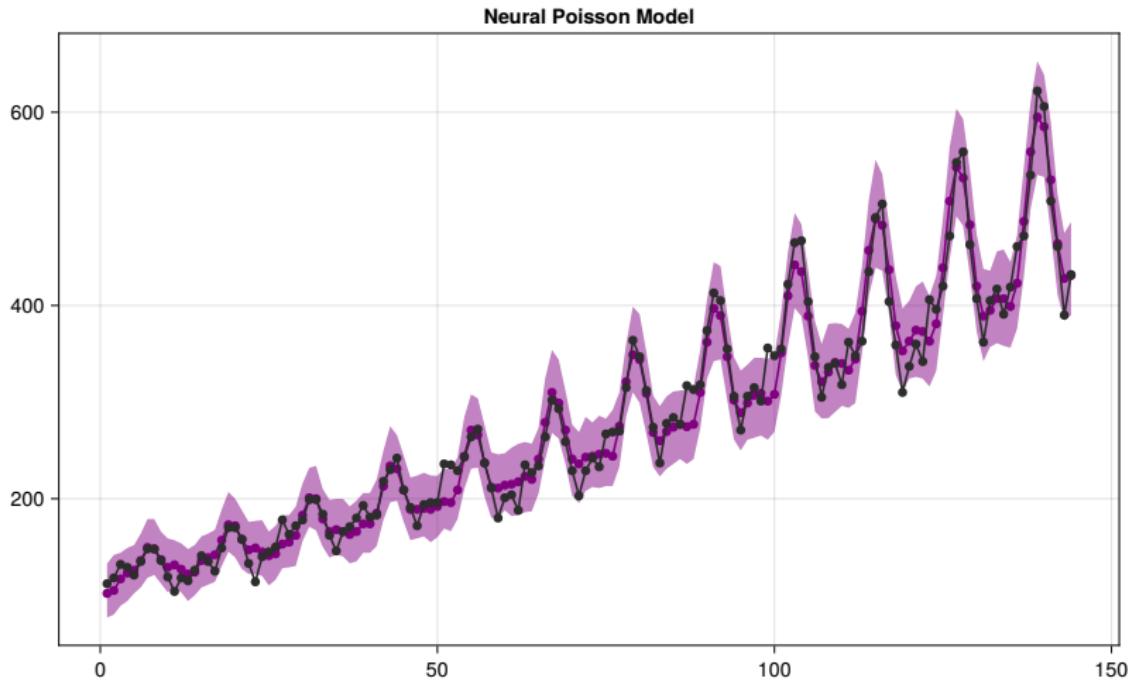
# Neural Poisson Regression

```
@model function neural_poisson_model(y)
    w ~ MvNormal(zeros(12), I)
    L1 = reshape(w[1:8], (4,2))
    L2 = reshape(w[9:12], (1,4))
    b = w[12:12]

    φ ~ Normal(0.0, 2.0)
    ω ~ Normal(0.0, 2.0)

    for t in 1:length(y)
        x = [t, cos(ω*(t-φ))]
        log_λ = ANN(L1, L2, b, x)[1]
        y[t] ~ Poisson(exp(log_λ))
    end
end
```

# Neural Poisson Regression



# Neural Poisson Regression

Of course, we don't have to manually define neural networks

[lux.jl Docs](#) |  Search | Ctrl K | [Home](#) | [Getting Started](#) | [Benchmarks](#) | [Tutorials](#) | [Manual](#) | [API](#) | [Version](#) |  |     0.6k ★ |

**LuxDL Docs**  
Elegant & Performant  
Scientific Machine  
Learning in JuliaLang

A Pure Julia Deep Learning Framework designed for  
Scientific Machine Learning



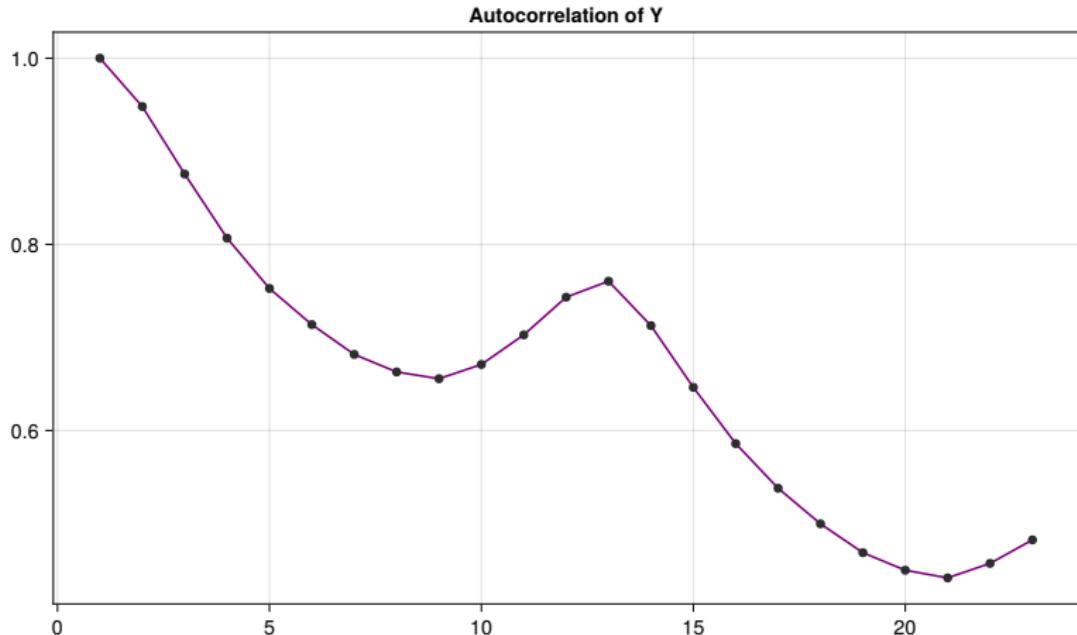
[Tutorials](#)

[API Reference](#) 

[View on GitHub](#)

# Autoregressive Models

A popular assumption in modeling time series data is **Autoregression**



# Autoregressive Models

$$Y_\tau \sim \text{Poisson}(\lambda_\tau)$$

$$\log \lambda_\tau = \alpha \log \lambda_{\tau-1} + \beta$$

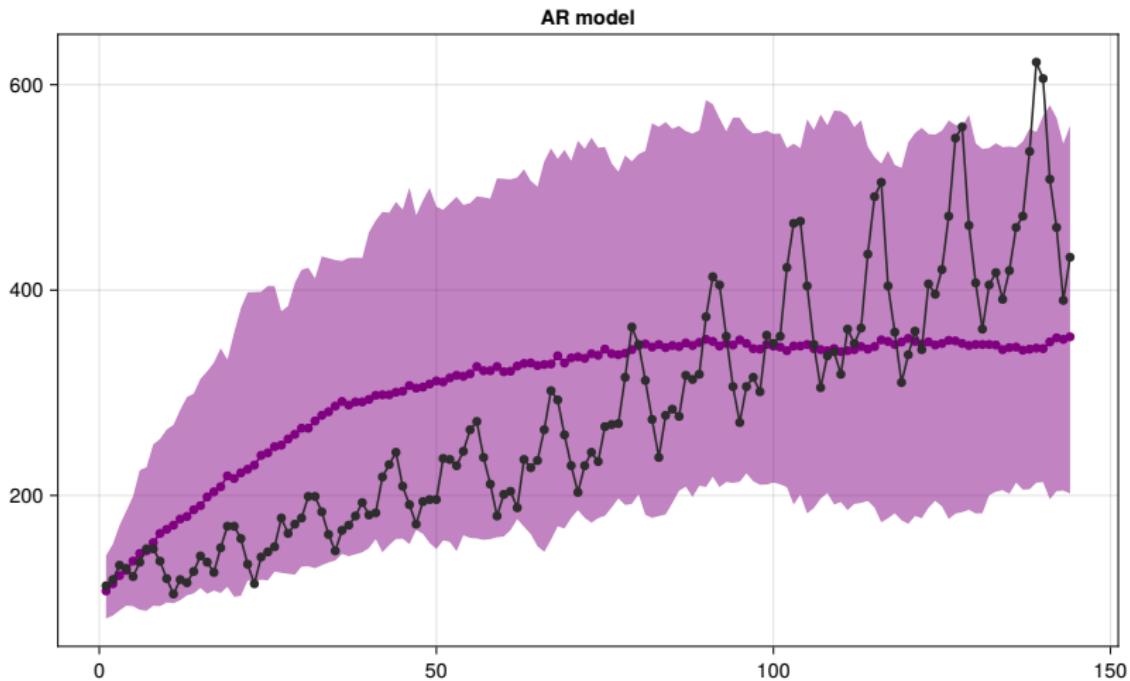
$$\alpha, \beta \sim \text{Normal}(0, 1)$$

$$\lambda_0 \sim \text{Poisson}(100.0)$$

```
@model function AR(y)
    α ~ Normal(0.0, 2.0)
    β ~ Normal(0.0, 2.0)
    y₀ ~ filldist(Poisson(100.0), 1)

    for t in 1:length(y)
        log_λ = α * log(lag(y, y₀, t-1)) + β
        λ = clamp(exp(log_λ), 100.0, 1000.0)
        y[t] ~ Poisson(λ)
    end
end
```

# Autoregressive Models



# Autoregressive Models

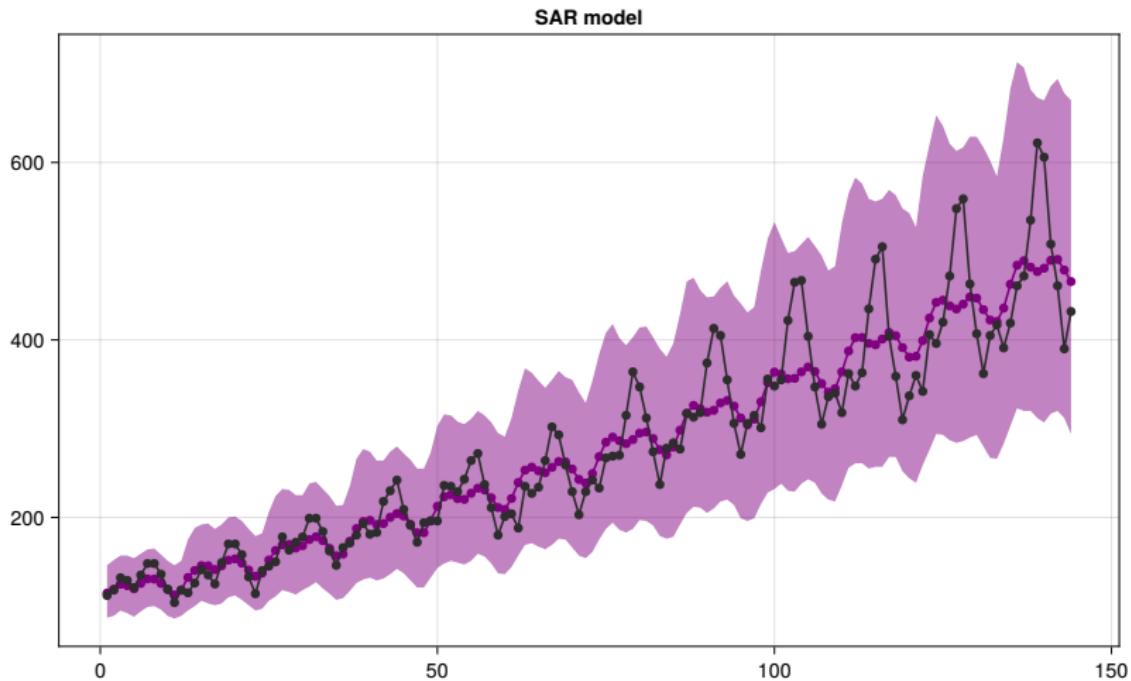
$$Y_\tau \sim \text{Poisson}(\lambda_\tau)$$

$$\log \lambda_\tau = \alpha_1 \log \lambda_{\tau-1} + \alpha_2 \log \lambda_{\tau-2} + \alpha_3 \log \lambda_{\tau-12} + \beta$$

```
@model function SAR(y)
    α1 ~ Normal(0.0, 2.0)
    α2 ~ Normal(0.0, 2.0)
    α3 ~ Normal(0.0, 2.0)
    β ~ Normal(0.0, 2.0)
    y0 ~ filldist(Poisson(100.0), 12)

    for t in 1:length(y)
        log_λ = α1 * log(lag(y, y0, t-1)) +
                α2 * log(lag(y, y0, t-2)) +
                α3 * log(lag(y, y0, t-12)) + β
        λ = clamp(exp(log_λ), 100.0, 1000.0)
        y[t] ~ Poisson(λ)
    end
end
```

# Autoregressive Models



# Autoregressive Models

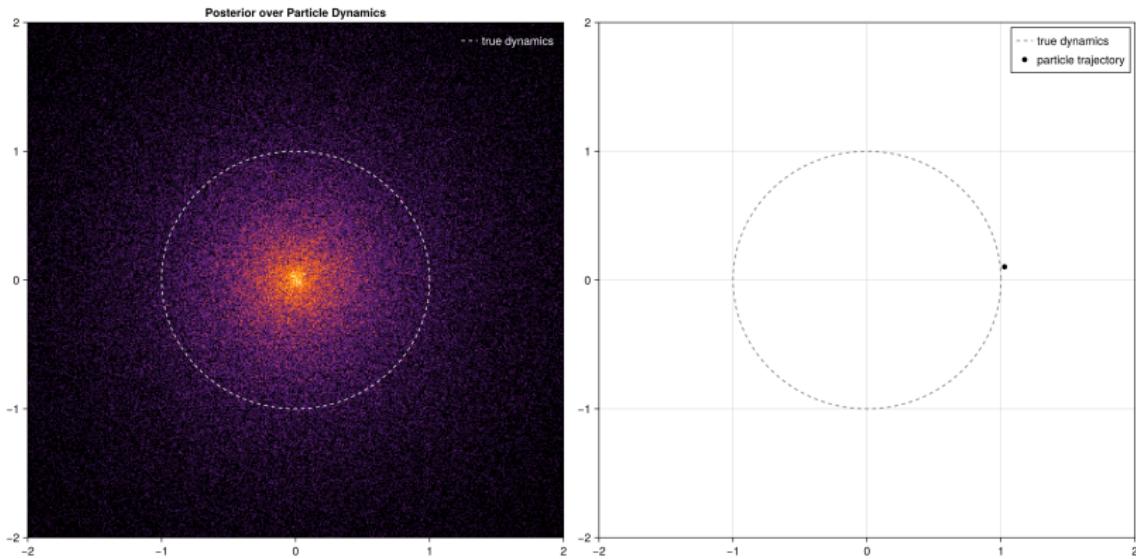
Generalizing to a multivariate scenarios is straightforward:

$$X_t \sim \text{MvNormal}(AX_{t-1}, \sigma I)$$
$$A \sim \text{MvNormal}(0, I)$$

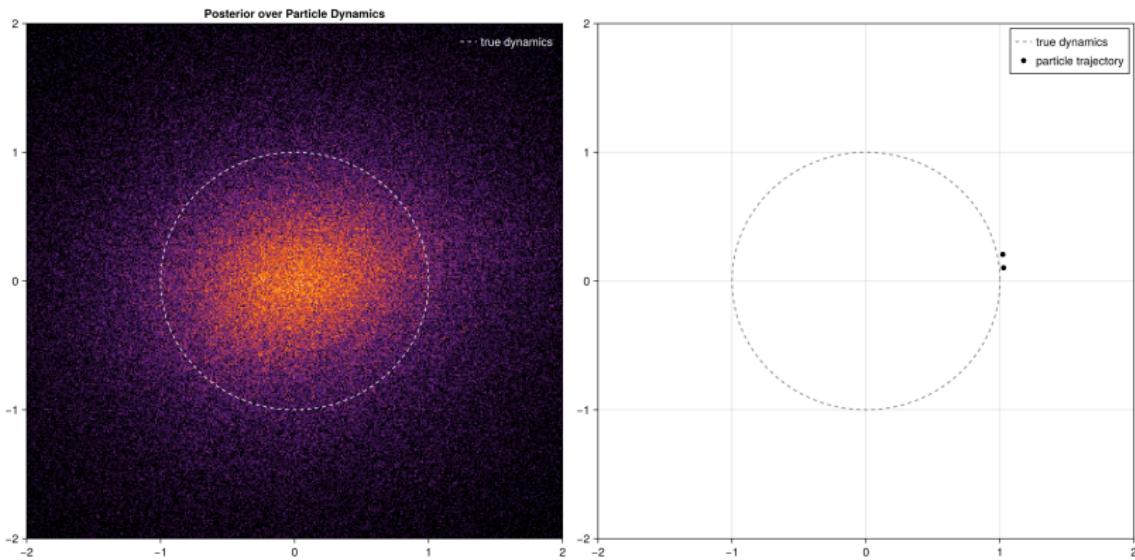
```
@model function VAR(x)
    A ~ MvNormal(zeros(4), I)
    σ ~ Exponential(1.0)

    for t in 2:length(x)
        x[t] ~ MvNormal(reshape(A, (2,2)) * x[t-1], σ * I(2))
    end
end
```

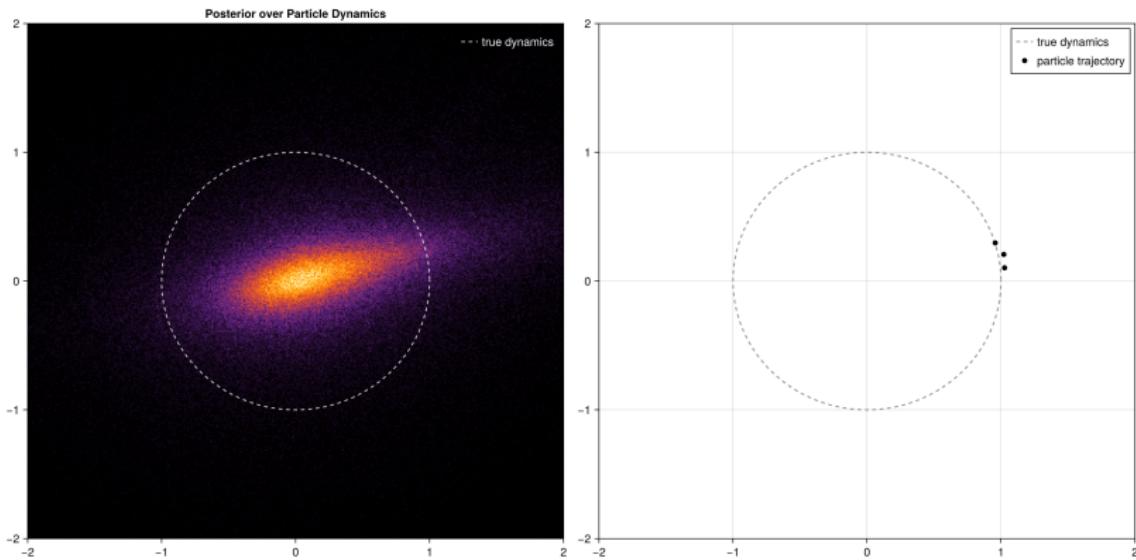
# Autoregressive Models



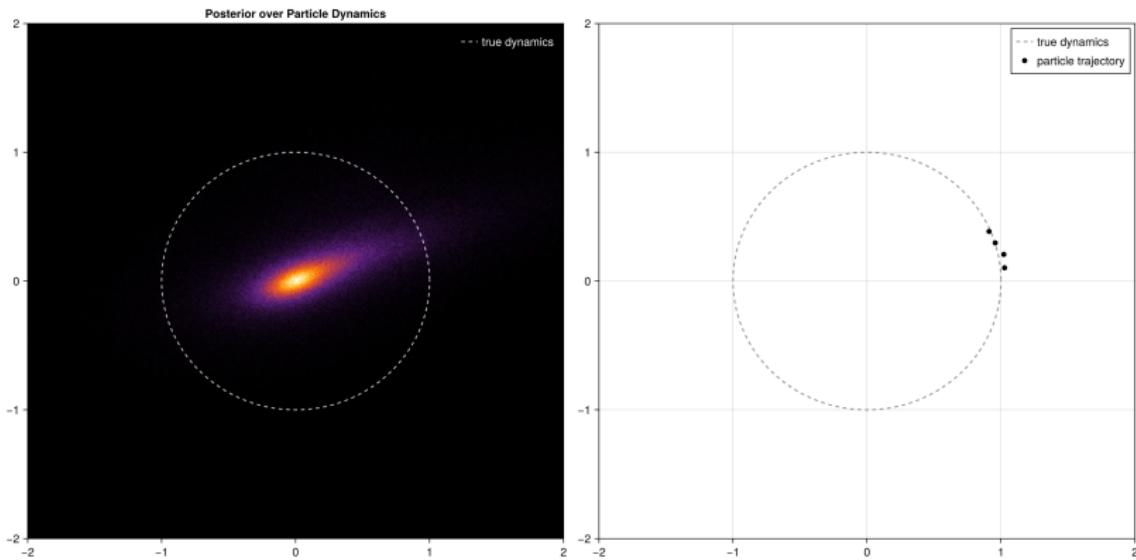
# Autoregressive Models



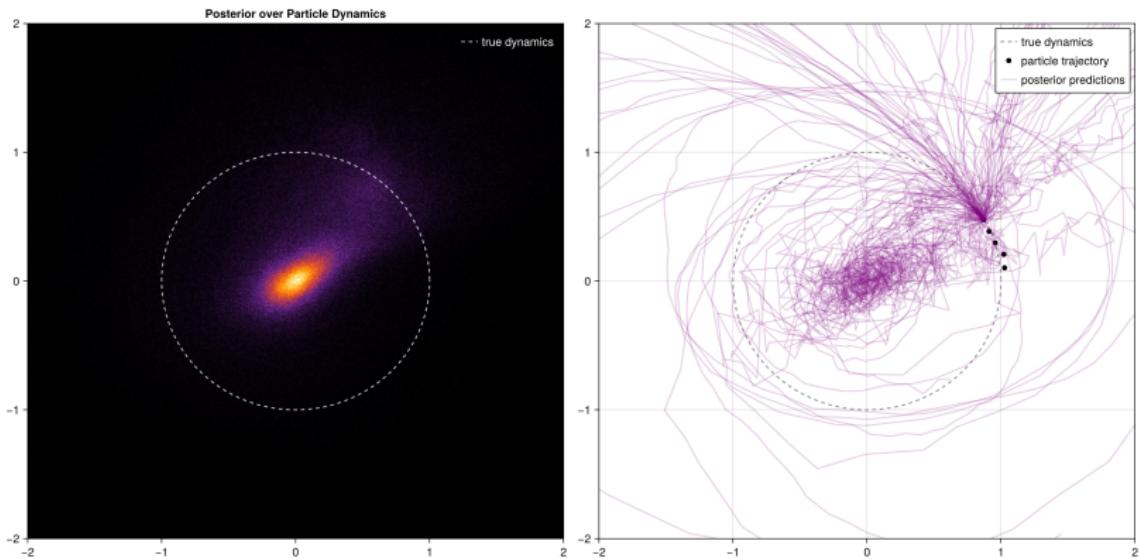
# Autoregressive Models



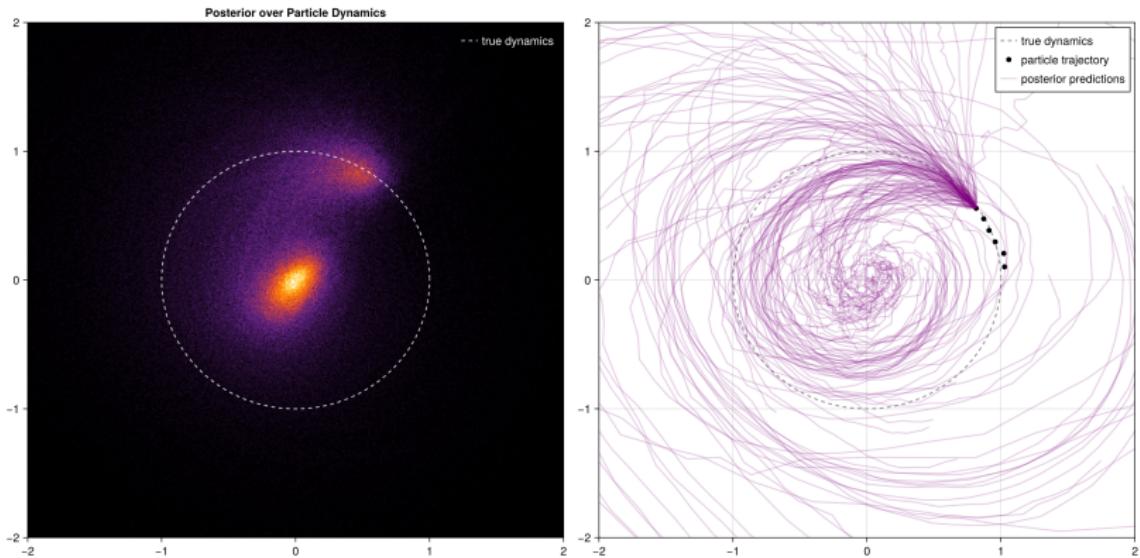
# Autoregressive Models



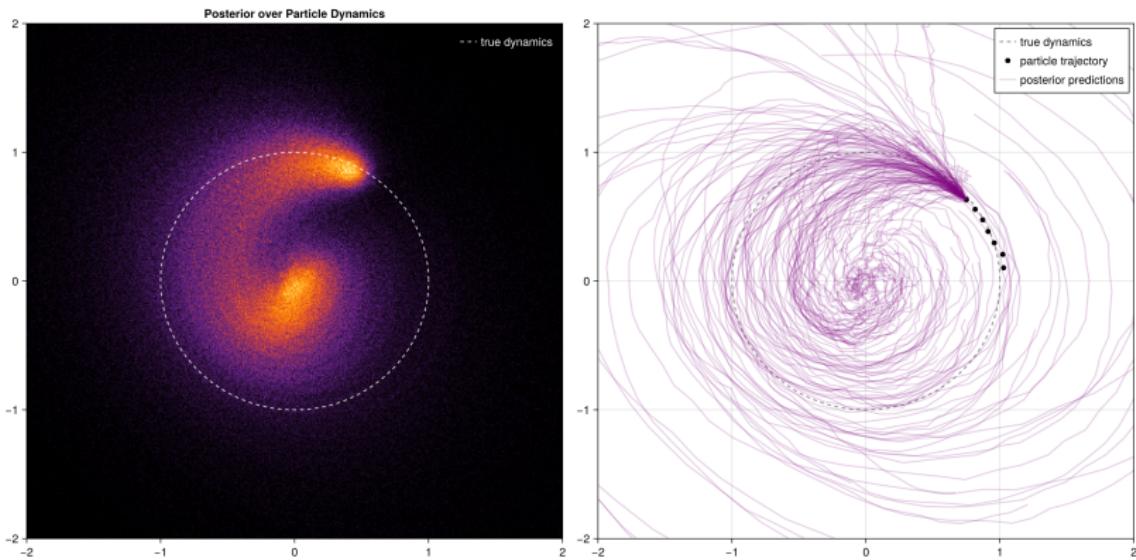
# Autoregressive Models



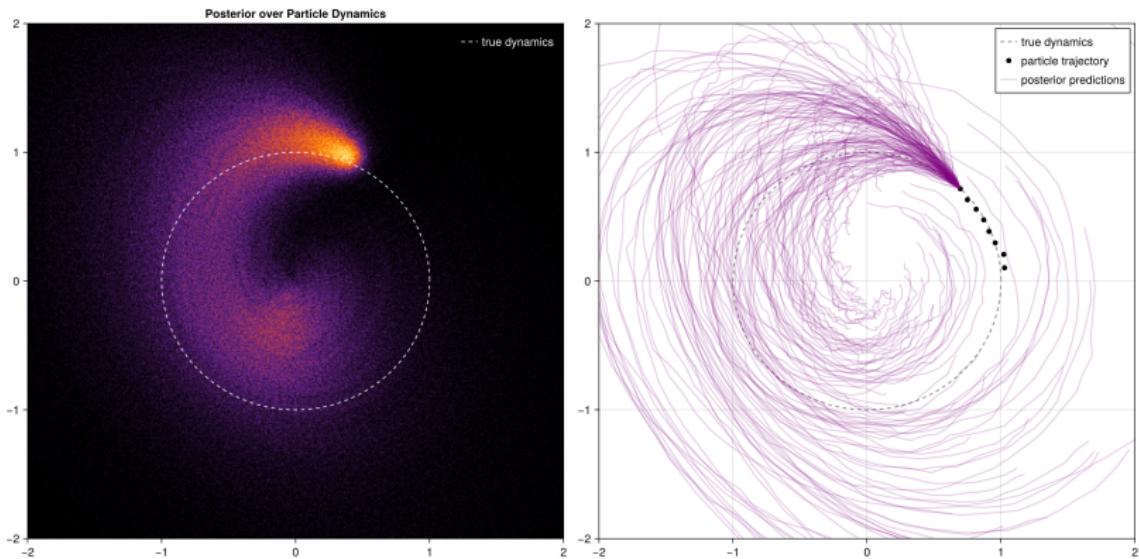
# Autoregressive Models



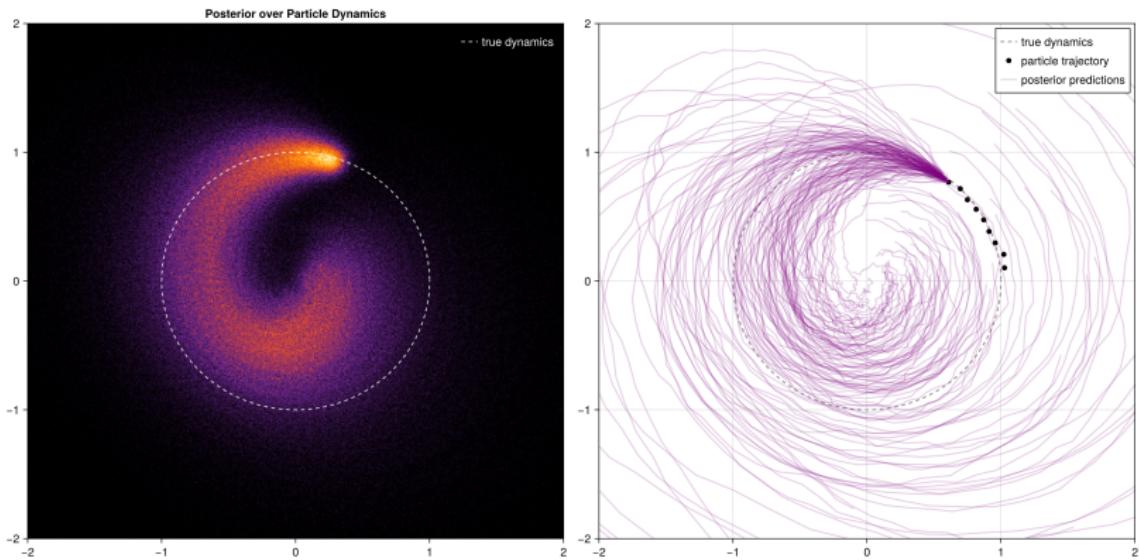
# Autoregressive Models



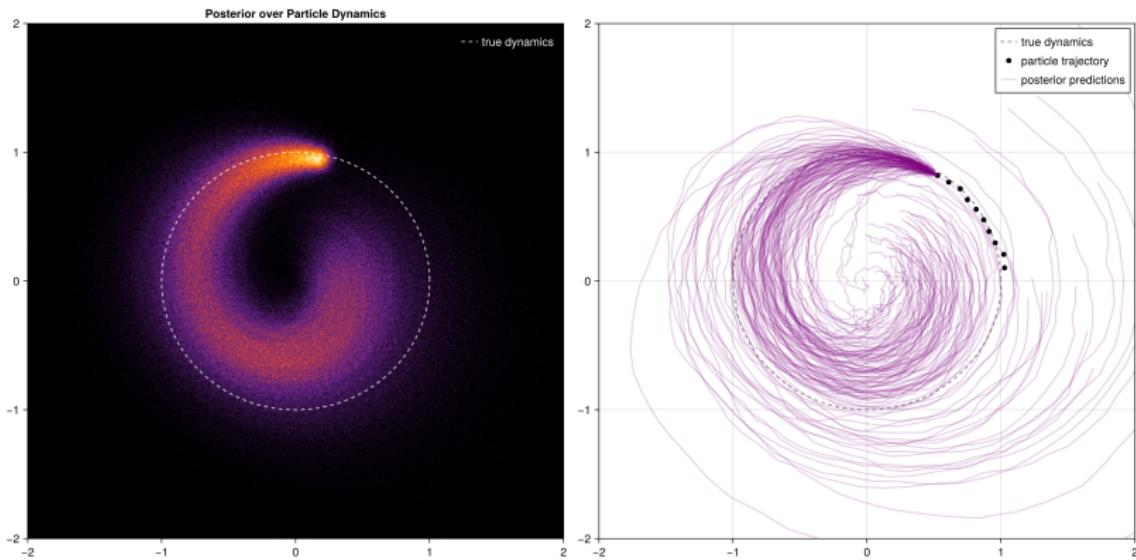
# Autoregressive Models



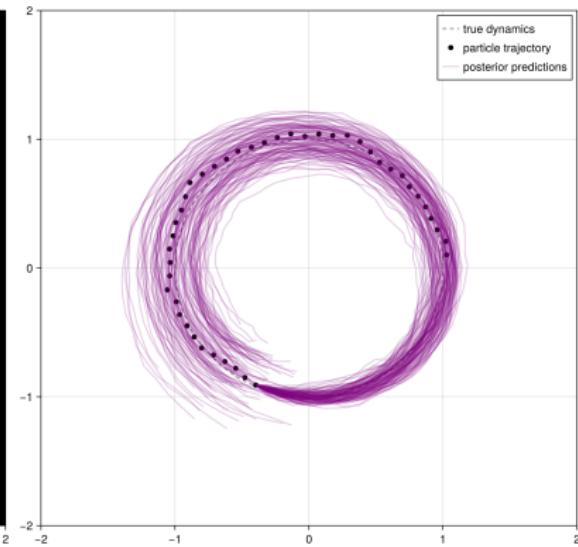
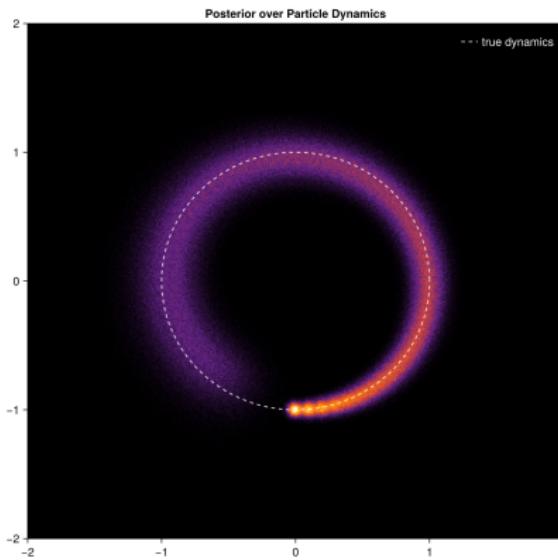
# Autoregressive Models



# Autoregressive Models



# Autoregressive Models



# Latent ODE

Autoregressive models and in general Dynamic Bayesian Networks are discrete-time models.

Can we find a continuous-time model which is computationally tractable?

## **Bayesian Differential Equations**

## Latent ODE

Remember our linear Poisson model  $\log \lambda_\tau = \alpha\tau + \beta$ . A similar model can be described by the following ODE with an unknown parameter and initial value:

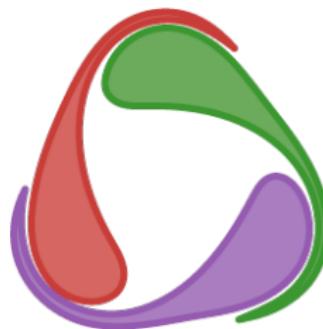
$$\frac{d \log \lambda}{d\tau} = \alpha$$

$$\log \lambda(1) = \beta$$

$$\alpha, \beta \sim \text{Normal}(0, 1)$$

We can solve ODEs in a probabilistic program and infer the unknown parameters!

# Latent ODE



SciML: Open Source Software for Scientific Machine Learning

- Advanced differential equation solvers
- Differentiable ODE solvers compatible with Turing

# Latent ODE

Install DifferentialEquations.jl:

```
Pkg.add("DifferentialEquations.jl")
```

Define the rate dynamics:

```
function rate_dynamics(dlogλ, logλ, p, t)
    α = p[1]
    dlogλ[1] = α
end
```

# Latent ODE

Now we can use the ODE in a probabilistic program:

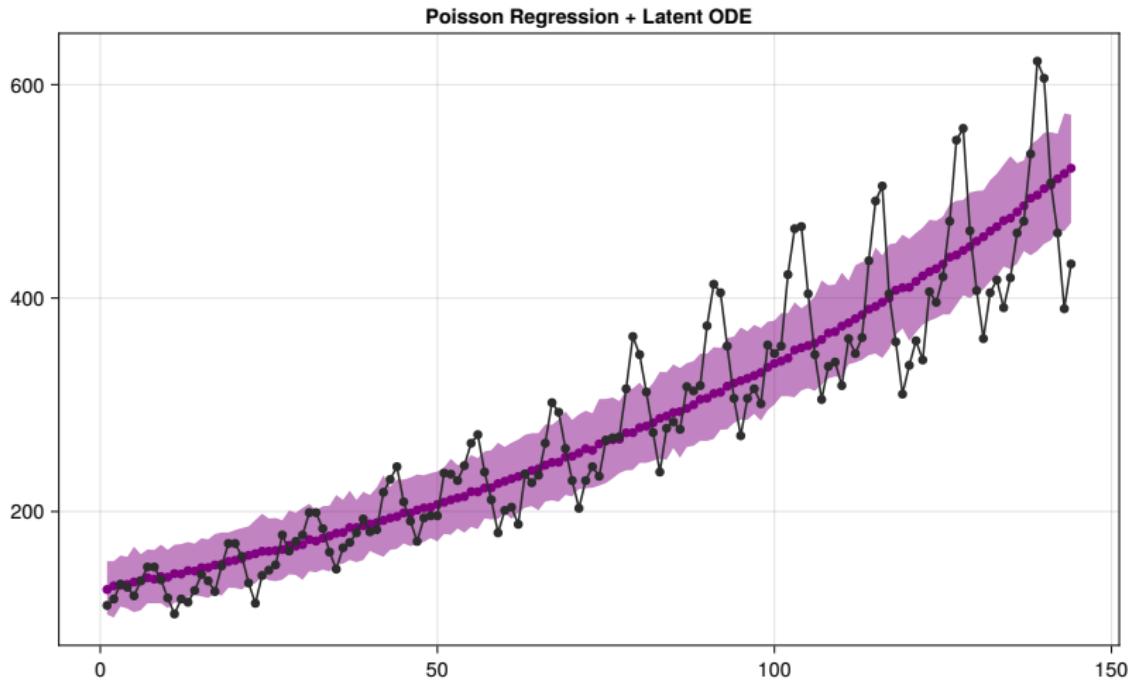
```
@model function latent_ODE_poisson(y)
    α ~ Normal(0.0, 2.0)
    β ~ Normal(0.0, 2.0)

    # solve the ODE for λ
    logλ₀ = [β]
    tspan = (1.0, length(y))
    p = [α]
    prob = ODEProblem(rate_dynamics, logλ₀, tspan, p)

    logλ = solve(prob, Tsit5(); saveat=1.0)

    for τ in 1:length(y)
        λ = clamp(exp(logλ[τ][1]), 100.0, 1000.0)
        y[τ] ~ Poisson(λ)
    end
end
```

# Latent ODE



# Latent ODE

Next we try a system of ODEs with latent parameters

$$\begin{aligned}\frac{du}{d\tau} &= \alpha_1 \\ \frac{dz_1}{d\tau} &= \alpha_2 z_2(\tau) + \alpha_3 \\ \frac{dz_2}{d\tau} &= \alpha_4 z_1(\tau) + \alpha_5 \\ \lambda(\tau) &= e^{u(\tau)} + e^{z_1(\tau)}\end{aligned}$$

```
function rate_dynamics(dlogλ, logλ, p, t)
    α1, α2, α3, α4, α5 = p
    dlogλ[1] = α1
    dlogλ[2] = α2 * logλ[3] + α3
    dlogλ[3] = α4 * logλ[2] + α5
end
```

# Latent ODE

